

Introduction to AI

SEARCH

Pedro Meseguer

Institut d'Investigació en Intel·ligència Artificial (IIIA)
Consejo Superior de Investigaciones Científicas (CSIC)



1

Contents

1. Introduction
2. Off-line search
 - ▶ Systematic search: formalization, A*, heuristics
 - ▶ Local search: several metaphors
3. On-line search
 - ▶ Single agent: unknown terrain
 - ▶ Two agents: adversarial search, Chess, Go
4. Search names
5. Wrap-up

2

2

1.Introduction

Example: 8-puzzle

1	5	2
4		3
6	7	8

3

1. 8 puzzle

- ▶ It is a 3x3 board with 8 numbered tiles and a hole (= blank space).
- ▶ A tile adjacent to the blank space can slide into the blank space.
- ▶ The objective is to reach the goal state.

1	2	3
8		6
7	5	4

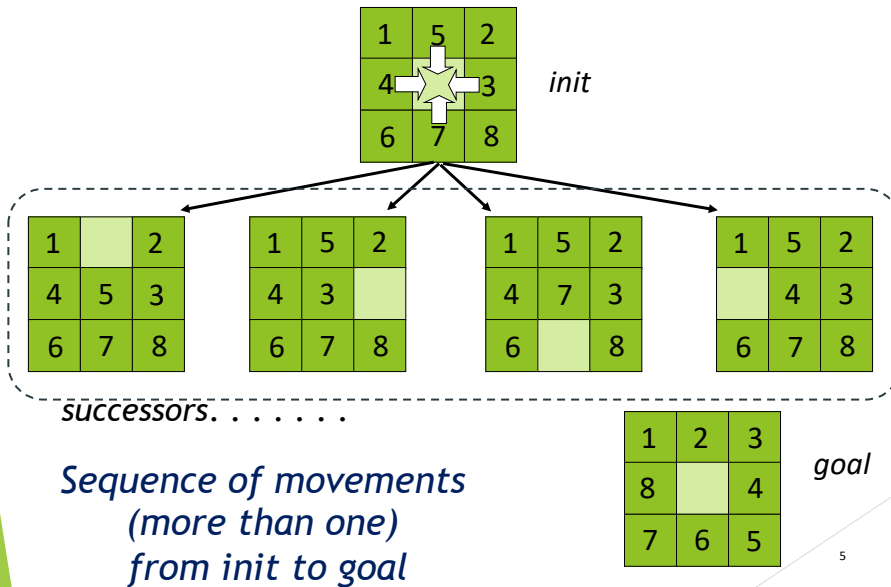
*initial
state*

1	2	3
8		4
7	6	5

*goal
state*

4

1.Example: Path finding in the 8 puzzle



5

1.Off-line search

Two separate phases:

1. Solution computation
2. Solution execution

Search for a complete solution

Execution

AI [search] is only concerned with the first phase

- ▶ Solution execution does not affect that solution:
in general, this does not happen in the real world
- ▶ For problems in very controlled environments

6

2. Concepts

- ▶ State: a possible problem configuration
- ▶ State space: all possible configurations (directed graph)
- ▶ Operators:
 - ▶ legal actions
 - ▶ they generate successors of a state
- ▶ States: initial and goal
explicit may be implicit
- ▶ Solution:
 - ▶ sequence operators initial to goal
(*sometimes*) goal state
- ▶ Problem instance: state-space plus initial and goal states

7


2. Example

1	5	2
4	3	
6	7	8

- ▶ State:
- ▶ State space: directed graph; 9! nodes; arcs are actions
- ▶ Operators:
 - ▶ legal actions are move the blank up, down, right, left
 - ▶ they generate successors of a state
- ▶ States: initial and goal
explicit may be implicit (test)
- ▶ Solution:
 - ▶ sequence of operators from initial to goal
- ▶ Problem instance: well defined

8

2.Example: Path finding in the 8 puzzle



9

2.Off-line systematic search: tree search

Systematic traversal of the state-space:

- ▶ It guarantees to find a solution, if one exists.
- ▶ Rubik's cube / Sliding puzzles: 8-puzzle, 15-puzzle, 24-puzzle


Exploring the state space of a problema as a tree:

- ▶ Root: initial state
- ▶ 1st level: the sucessors of root
- ▶ 2nd level: the sucessors of sucessors of the root
- ▶ ...
- ▶ dth level: the sucessors of nodes at level d-1

10

2. Off-line systematic search: blind search

Search tree:



Depth-first search (DFS): explore by **branches**

Breadth-first search (BFS): explore by **levels**

11

2. Off-line systematic search: A* algorithm

A* algorithm:

- ▶ important in systematic search: best-first heuristic search
- ▶ known from long time (since 1968)

A* expand nodes (=generate successors) with minimum f

- ▶ node lists: open (to be expanded) and closed (already expanded)
- ▶ where $f(x) = g(x) + h(x)$
 - ▶ $g(x)$: the cost already spent to reach x from init
 - ▶ $h(x)$: the expected cost to reach a solution from x (*heuristic*)
- ▶ expand the node with minimum f in open list : gen. its successors

12

2. Off-line systematic search: heuristics

Heuristic:

- ▶ Estimates the distance to the closest goal
- ▶ Admissibility: never overestimates the real distance
- ▶ A* with an admissible heuristic: an optimal path to the goal
- ▶ Cheap to compute (impact in practice)

Examples of heuristics:

- ▶ 8-puzzle:
 - ▶ #tiles out of place
 - ▶ Sum Manhattan distance for each tile out of place

13

2. Off-line systematic search: combinatorial explosion

Size of state space grows quickly as problem size increases

▶ Sliding puzzles:

- ▶ 8-puzzle 9!
- ▶ 15-puzzle 16!
- ▶ 24-puzzle 25!

The state space of these problems is divided in two, unconnected parts of the same size

- ▶ Main reason of the failure of early AI

Search (weak methods)

- ▶ are overwhelmed by huge state spaces
- ▶ only work for *toy problems* (memo causing AI winter 1966)

14

2. Off-line local search

- ▶ Optimization: min cost function; solution: global minimum
- ▶ May fail finding a solution, even if it exists
- ▶ Simplest: Hill-climbing
- ▶ Often inspired by natural processes:
 - ▶ simulated annealing
 - ▶ genetic algorithms
 - ▶ ants algorithms
- ▶ Basically,
 - ▶ they do not keep track of states previously visited
 - ▶ look forward for immediate improvement (greediness)

15

2. Off-line local search: hill-climbing

Hill-climbing: minimize $f(x)$

- ▶ Generate successors of current state
- ▶ From x move to y such that $f(x) > f(y)$: the move improves cost function
- ▶ Stochastic search

Local optima:

- ▶ $f(x) \leq f(w)$, for all w in $\text{succ}(x)$
- ▶ how escape from local minima?

Escape strategies:

- ▶ Accept moves to states of higher function values with low probability (simulated annealing).


16

16

2. Off-line local search: genetic algorithms

Knapsack problem:

Given a set of items, each with a weight and a value, determine the number of each item to include in a collection so that the total weight is less than or equal to a given limit and the total value is as large as possible.



17

Genetic algorithms: Chromosome / fitness function

A candidate solution is a vector = chromosome

1	0	1	0	0
---	---	---	---	---



To discriminate among chromosomes:

fitness function $F(V)$: total value of vector V

$$F(\begin{array}{|c|c|c|c|c|} \hline 1 & 0 & 1 & 0 & 0 \\ \hline \end{array}) = \$2 + \$10 = \$12$$

18

2. Off-line local search: genetic algorithms

- ▶ Parallel search: a pool of solutions randomly generated (strings)
- ▶ Combines: exploration & exploitation
- ▶ Applying operators: generation(i) \rightarrow offspring, generation (i+1)
- ▶ Operators:
 - ▶ Selection : two strings are selected
 - ▶ Crossover: old strings are cut \rightarrow new strings are obtained
 - ▶ Mutation: some bits may change randomly

Iterative process:


- ▶ Several generations are produced, each with better individuals
- ▶ Until a solution is found or exhaust resources

19

Genetic algorithms: Selection


For reproduction we select:

- The best chromosomes, according to the fitness function




20

Genetic algorithms: Reproduction (crossover)



21

Genetic algorithms: Mutation



21/9/22

22


22

Genetic algorithms: New generation

Adjust the selection, crossover and mutation to obtain a new generation of K chromosomes.

The new generation replaces the old one.

The whole process repeats.



23

2. Off-line local search: suboptimality

When finding a solution:

- ▶ no guarantee that the path is optimal
(= minimum cost)

Often

- ▶ good quality solutions
- ▶ efficiency achieved

24

3. On-line search

► Off-line search:

Search for a complete solution

Execution

► unrealistic in many domains

► On-line search: interleave

- Solution computation
- Solution execution

► Domains:

- Unknown terrain
- Dynamic

► Families:

- Single-agent
- Two agents in games, adversarial search

25

3 On-line search: single-agent

Age of Empires



Warcraft III



26

26

3. On-line search: single-agent approaches


- ▶ Two families:
 - ▶ Incremental search
 - ▶ Real-time search
- ▶ Basic difference: time for first move
- ▶ If time for first move could be high
 - ▶ Incremental search
- ▶ Otherwise
 - ▶ Real-time search

27

27

On-line search: Path-finding in unknown terrain


How?



28

28

Off-line Search: A*



29

On-line search: Incremental

- ▶ As classical heuristic (off-line) search
- ▶ Suitable for **unknown environments**
- ▶ On-line search:

Search for a complete solution from the start state

Execution


If unfeasible, search for a new complete solution from the current state

Execution

... iterates until finding a goal

30

On-line search: incremental, D*



31


Real-Time Heuristic Search

- ▶ Search on a local space around the current state
- ▶ As result of search:
 - ▶ Heuristic of some states are updated (*learning space*)
 - ▶ Move(s) in the local search space
- ▶ Solution is no optimal:
 - ▶ approaches converge to optimality after repeated exec.
- ▶ Interleave search and action execution
- ▶ On-line search:



32

Real-time Search: LRTA*



33

3. On-line search. Two agents, zero sum, perfect information games: adversarial search



34

34

3. On-line search. Two agents, zero sum, perfect information games: adversarial search

- ▶ Two players: A and B
- ▶ Perfect information:
 - ▶ Each player knows all the information of the opponent
 - ▶ No random elements
 - ▶ *Chess, checkers, otello, go*
- ▶ Excluded:
 - ▶ Incomplete information games: *poker, bridge*
 - ▶ Stochastic games (dices): *backgammon*
- ▶ Zero-sum: if utility for A is $x \Rightarrow$ utility for B is $-x$


35

Game Programs

- ▶ Programs world championship: *checkers, othello*
- ▶ Program with good performance: *chess, go*
 - ▶ important: since 1956, *chess* was an IA goal
 - ▶ *Deep Blue* won *Kasparov* in 1997
 - ▶ *AlphaGo* won *Lee Sedol* in 2015
- ▶ Clear economic importance
- ▶ Strategies:
 - ▶ *brute-force*: search in spaces of billions of nodes
 - ▶ *smart* sampling, if the problem is too large

36


Game Tree



- Players alternate by levels
- Node successors: all legal moves that the current player can do

37

Example: Chess



38

MiniMax Main Idea

Search and back-propagate the value of the terminals
(from Alan Turing, late 40's) [Turing chess program](#)

Terminal nodes: value 1, -1, 0 (A wins, B wins, draw)

Two types of nodes: **max** wins with 1 / **min** wins with -1


Back-propagation:

- ▶ **max**: the maximum of the values of the children
- ▶ **min**: the minimum of the values of the children


39

MiniMax Game Tree

1 win ■
-1 win ●
0 draw



What is the best move? Back-propagate values from node terminals, assuming that each player selects the most favourable move for him/her



EAIA

40


3. On-line adversarial search: Go



41

However, for *Go* this approach does not work!

- ▶ *Go*:
 - ▶ complex (branching~250), large (depth~150), very dynamic (no quiescence)
 - ▶ no good evaluation function
- ▶ Brute-force approaches do not work
- ▶ What about some kind of sampling? → Monte Carlo Tree Search (from 2005 on)

4
2


42

Monte Carlo Tree Search (MCTS)

Basic idea:

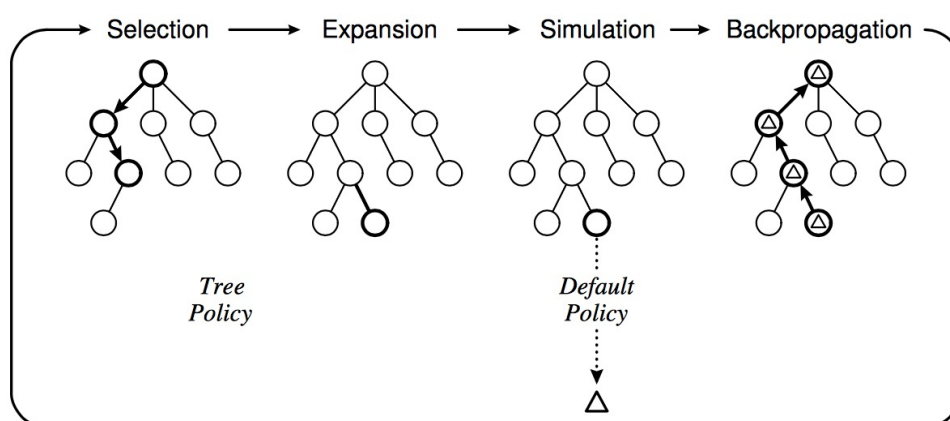
- ▶ keep (develop and store) a small part of the game tree close to the root: partial tree
- ▶ from the fringe of that partial tree, run a simulation until a terminal node
- ▶ update node statistics in the partial tree

Iterate until some limit (time or #iterations); then select the *best* move



43

MCTS: Phases



44

4.Search names: Nils Nilsson



- ▶ (1933 – 2019)
- ▶ Full profesor Stanford
- ▶ One of the big old names
- ▶ Among the creators of the A* algorithm (1968)
- ▶ Shakey robot -> STRIPS planner (1971) [*planning*]

45

4.Search names: Judea Pearl



- ▶ (1936 –)
- ▶ Full professor UCLA (retired)
- ▶ Known by his books: *Heuristics*, 1984; *Probabilistic Reasoning*, 1988; *Causality*, 2000
- ▶ Contributions to: Bayesian networks [*uncertainty*]

46

4.Search names: Richard Korf



- ▶ (195? –)
- ▶ Full profesor UCLA
- ▶ Many contributions to systematic search:
 - ▶ Off-line search: frontier search
 - ▶ On-line search: real-time heuristic search

47

4.Search names: Sven Koenig



- ▶ (195? –) German
- ▶ Full profesor USC (Southern California)
- ▶ Many contributions to single-agent search:
 - ▶ Partially known, non-stationary, non-deterministic domains
 - ▶ [*multi agent planning, robotics, videogames*]

48

4. Search names: Deep Blue

- ▶ 1997, program developed by IBM
- ▶ Mastering chess, won to the world champion Gary Kasparov
- ▶ Combination of:
 - ▶ Parallel alpha-beta search
 - ▶ Variable depth (Singular extensions)
 - ▶ Library of openings and endings



49

49

4. Search names: AlphaGo



AlphaGo

- ▶ 2015, program developed by Google DeepMind
- ▶ Mastering the Go game (more difficult than chess), won a 5-games match against the unofficial world champion Lee Sedol
- ▶ Combination of
 - ▶ Monte-Carlo Tree Search
 - ▶ Deep neural networks

50

50

5.Wrap-up (I)

- ▶ Basics:
 - ▶ A state
 - ▶ State-space of a problem
 - ▶ Successors of a state
 - ▶ Path-finding
- ▶ Off-line search
 - ▶ Systematic search
 - ▶ Local search:
 - ▶ hill climbing [local optima, global optimum]
 - ▶ genetic algorithms:
 - ▶ chromosome, fitness function
 - ▶ selection, crossover, mutation

51

5.Wrap-up (II)

- ▶ On-line search
 - ▶ Single agent search
 - ▶ Incremental search
 - ▶ Real-time search
 - ▶ Two-agent, adversarial search
 - ▶ Game tree
 - ▶ Games: chess, go
- ▶ Search names: Nilsson, Pearl, Korf, Koenig, Deep Blue, Alpha Go
- ▶ Algorithms mentioned: DFS, BFS, A*, genetic, MiniMax
- ▶ Examples: 8-puzzle, knapsack

52

Further reading

- ▶ Russel & Norvig 3rd ed:
3.1, 3.2, 3.3, 4.1, 5.1, 5.2, 5.7 plus
Bibliographical Notes chapters 3, 4 and 5
- ▶ Heuristics (Pearl, 1984)
- ▶ Korf's talk at ICAPS 2018 (very instructive, no local search)
<https://www.youtube.com/watch?v=X6qCBcubZIE>

53