

## 2.Off-line local search

- Optimization: min cost function; solution: global minimum
- May fail finding a solution, even if it exists
- Simplest: Hill-climbing
- Often inspired by natural processes:
  - simulated annealing
  - genetic algorithms
  - ants algorithms
- Basically,
  - they do not keep track of states previously visited
  - look forward for immediate improvement (greediness)

1

## 2.Off-line local search: hill-climbing

Hill-climbing: minimize  $f(x)$

- Generate successors of current state
- From  $x$  move to  $y$  such that  $f(x) > f(y)$ : the move that decreases most the cost function
- Stochastic search

Local optima:

- $f(x) \leq f(w)$ , for all  $w$  in  $\text{succ}(x)$
- how escape from local minima?

Escape strategies:


- Accept moves to states of higher function values with low probability (simulated annealing).

2

## 2.Off-line local search: genetic algorithms

Knapsack problem:

Given a set of items, each with a weight and a value, determine the number of each item to include in a collection so that the total weight is less than or equal to a given limit and the total value is as large as possible.



3

## Genetic algorithms: Chromosome / fitness function

A candidate solution is a vector = chromosome

1 0 1 0 0



To discriminate among chromosomes:

*fitness function*  $F(V)$ : total value of vector  $V$

$$F(1 \ 0 \ 1 \ 0 \ 0) = \$2 + \$10 = \$12$$

4

## 2. Off-line local search: genetic algorithms

- Parallel search: a pool of solutions randomly generated (strings)
- Combines: exploration & exploitation
- Applying operators:  $\text{generation}(i) \rightarrow \text{offspring}, \text{generation}(i+1)$
- Operators:
  - Selection : two strings are selected
  - Crossover: old strings are cut  $\rightarrow$  new strings are obtained
  - Mutation: some bits may change randomly

Iterative process:


- Several generations are produced, each with better individuals
- Until a solution is found or exhaust resources

5

## Genetic algorithms: Selection


For reproduction we select:

- The best chromosomes, according to the fitness function




6

## Genetic algorithms: Reproduction (crossover)



7

## Genetic algorithms: Mutation




8

## Genetic algorithms: New generation

Adjust the selection, crossover and mutation to obtain a new generation of  $K$  chromosomes.

The new generation replaces the old one.

The whole process repeats.



9

## 2. Off-line local search: suboptimality

When finding a solution:

- ▶ no guarantee that the path is optimal (= minimum cost)

Often

- ▶ good quality solutions
- ▶ efficiency achieved

10

## 3. On-line search

▶ Off-line search:

Search for a complete solution

Execution

▶ unrealistic in many domains

▶ On-line search: interleave

- ▶ Solution computation
- ▶ Solution execution

▶ Domains:

- ▶ Unknown terrain
- ▶ Dynamic

▶ Families:

- ▶ Single-agent
- ▶ Two agents in games, adversarial search

11

## 3 On-line search: single-agent

Age of Empires

Warcraft III




12

### 3. On-line search: single-agent approaches

- ▶ Two families:
  - ▶ Incremental search
  - ▶ Real-time search
- ▶ Basic difference: time for first move
- ▶ If time for first move could be high
  - ▶ Incremental search
- ▶ Otherwise
  - ▶ Real-time search

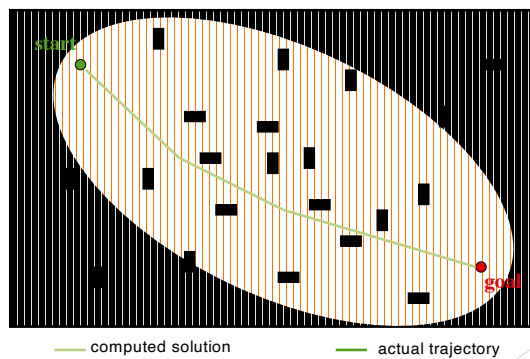
13

### On-line search: Path-finding in unknown terrain



14

### Off-line Search: A\*



15

### On-line search: Incremental

- ▶ As classical heuristic (off-line) search
- ▶ Suitable for **unknown environments**
- ▶ On-line search:


|   |           |
|---|-----------|
| Search for a complete solution from the start state | Execution |
|---|-----------|

|  |           |
|--|-----------|
| If unfeasible, search for a new complete solution from the current state | Execution |
|--|-----------|

... iterates until finding a goal

16

### On-line search: incremental, D\*



17

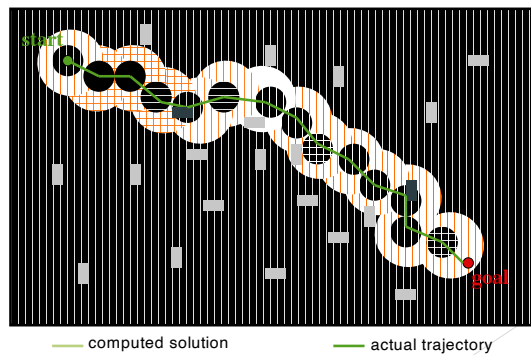
### Real-Time Heuristic Search

- ▶ Search on a local space around the current state
- ▶ As result of search:
  - ▶ Heuristic of some states are updated (*learning space*)
  - ▶ Move(s) in the local search space
- ▶ Solution is no optimal:
  - ▶ approaches converge to optimality after repeated exec.
- ▶ Interleave search and action execution
- ▶ On-line search:



18

### Real-time Search: LRTA\*



19

### 3. On-line search. Two agents, zero sum, perfect information games: adversarial search



20

### 3. On-line search. Two agents, zero sum, perfect information games: adversarial search

- ▶ Two players: A and B
- ▶ Perfect information:
  - ▶ Each player knows all the information of the opponent
  - ▶ No random elements
  - ▶ *Chess, checkers, otello, go*
- ▶ Excluded:
  - ▶ Incomplete information games: *poker, bridge*
  - ▶ Stochastic games (dices): *backgammon*
- ▶ Zero-sum: if utility for A is  $x \Rightarrow$  utility for B is  $-x$


21

### Game Programs

- ▶ Programs world championship: *checkers, othello*
- ▶ Program with good performance: *chess, go*
  - ▶ important: since 1956, *chess* was an IA goal
  - ▶ *Deep Blue* won *Kasparov* in 1997
  - ▶ *AlphaGo* won *Lee Sedol* in 2015
- ▶ Clear economic importance
- ▶ Strategies:
  - ▶ *brute-force*: search in spaces of billions of nodes
  - ▶ *smart sampling*, if the problem is too large

22


### Game Tree



- Players alternate by levels
- Node successors: all legal moves that the current player can do

23

### Example: Chess



24

### MiniMax Main Idea

Search and back-propagate the value of the terminals  
(from Alan Turing, late 40's) [Turing chess program](#)

Terminal nodes: value 1, -1, 0 (A wins, B wins, draw)

Two types of nodes: **max** wins with 1 / **min** wins with -1


Back-propagation:

- **max**: the maximum of the values of the children
- **min**: the minimum of the values of the children


25

### MiniMax Game Tree

1 win ■  
-1 win ●  
0 draw



What is the best move? Back-propagate values from node terminals, assuming that each player selects the most favourable move for him/her



26


### 3. On-line adversarial search: Go



27

### However, for *Go* this approach does not work!

- **Go**:
  - complex (branching-250), large (depth-150), very dynamic (no quiescence)
  - no good evaluation function
- Brute-force approaches do not work
- What about some kind of sampling? → Monte Carlo Tree Search (from 2005 on)




28

## Monte Carlo Tree Search (MCTS)

Basic idea:


- keep (develop and store) a small part of the game tree close to the root: partial tree
- from the fringe of that partial tree, run a simulation until a terminal node
- update node statistics in the partial tree

Iterate until some limit (time or #iterations); then select the *best move*



29

## MCTS: Phases



30

## 4.Search names: Nils Nilsson



- (1933 – 2019)
- Full professor Stanford
- One of the big old names
- Among the creators of the A\* algorithm (1968)
- Shakey robot -> STRIPS planner (1971) [planning]

31

## 4.Search names: Judea Pearl



- (1936 – )
- Full professor UCLA (retired)
- Known by his books: *Heuristics*, 1984; *Probabilistic Reasoning*, 1988; *Causality*, 2000
- Contributions to: Bayesian networks [uncertainty]

32



#### 4.Search names: Richard Korf



- ▶ (195? – )
- ▶ Full profesor UCLA
- ▶ Many contributions to systematic search:
  - ▶ Off-line search: frontier search
  - ▶ On-line search: real-time heuristic search

33

#### 4.Search names: Sven Koenig



- ▶ (195? – ) German
- ▶ Full profesor USC (Southern California)
- ▶ Many contributions to single-agent search:
  - ▶ Partially known, non-stationary, non-deterministic domains
  - ▶ *[multi agent planning, robotics, videogames]*

34

#### 4.Search names: Deep Blue

- ▶ 1997, program developed by IBM
- ▶ Mastering chess, won to the world champion Gary Kasparov
- ▶ Combination of:
  - ▶ Parallel alpha-beta search
  - ▶ Variable depth (Singular extensions)
  - ▶ Library of openings and endings



35

#### 4.Search names: AlphaGo



- ▶ 2015, program developed by Google DeepMind
- ▶ Mastering the Go game (more difficult than chess), won a 5-games match against the unofficial world champion Lee Sedol
- ▶ Combination of
  - ▶ Monte-Carlo Tree Search
  - ▶ Deep neural networks

36

## 5.Wrap-up (I)

- ▶ Basics:
  - ▶ A state
  - ▶ State-space of a problem
  - ▶ Successors of a state
  - ▶ Path-finding
- ▶ Off-line search
  - ▶ Systematic search
  - ▶ Local search:
    - ▶ hill climbing [local optima, global optimum]
    - ▶ genetic algorithms:
      - ▶ chromosome, fitness function
      - ▶ selection, crossover, mutation

37

## 5.Wrap-up (II)

- ▶ On-line search
  - ▶ Single agent search
    - ▶ Incremental search
    - ▶ Real-time search
  - ▶ Two-agent, adversarial search
    - ▶ Game tree
    - ▶ Games: chess, go
- ▶ Search names: Nilsson, Pearl, Korf, Koenig, Deep Blue, Alpha Go
- ▶ Algorithms mentioned: DFS, BFS, A\*, genetic, MiniMax
- ▶ Examples: 8-puzzle, knapsack

38

## Further reading

- ▶ Russel & Norvig 3rd ed:
  - 3.1, 3.2, 3.3, 4.1, 5.1, 5.2, 5.7 plus
  - Bibliographical Notes chapters 3, 4 and 5
- ▶ Heuristics (Pearl, 1984)
- ▶ Korf's talk at ICAPS 2018 (very instructive, no local search)
  - <https://www.youtube.com/watch?v=X6qCBcubZIE>

39